

Create the AWS Lambda Function

This Lambda function will be the core of our image processing pipeline. It will:

- Trigger on new image uploads: Specifically, when an image is placed in our input S3 bucket.
- Resize/compress: It will contain the logic to transform the image (in our case, resize it).
- Save the processed image: The output will be stored in our designated output S3 bucket.

Let's build it step-by-step:

Create the Lambda Function Skeleton

1. Go to the Lambda Console: Navigate to the AWS Lambda service in your console.
2. Initiate Function Creation: Click the "Create function" button.
3. Configure Basic Settings:
 - "Author from scratch": Select this option.
 - "Function name": Enter ImageProcessorFunction
 - "Runtime": Choose Python 3.12
 - "Permissions": This is where we link our IAM role!
 - Expand the "Change default execution role" section.
 - Choose "Use an existing role".
 - From the dropdown, select your previously created role: lambda-image-processing-role.
4. Create the Function: Click "Create function".

Add S3 Trigger to the Lambda Function

Now that the function exists, we need to tell it when to run. We'll configure an S3 trigger so it's invoked every time a new image lands in our input bucket.

1. Access Your Lambda Function: If you're not already there, go back to the Lambda Console and open your newly created function: ImageProcessorFunction.
2. Add a New Trigger:
 - In the "Function overview" section, click on "Add trigger".
3. Configure the S3 Trigger:
 - "Select a trigger": Choose "S3".
 - "Bucket": Select your input bucket from the dropdown: image-processing-input. (Make sure you pick the correct one!)
 - "Event type": Choose "All objects create events" (or specifically "Put" for new uploads).
 - "Prefix / Suffix": These are optional but highly recommended for efficiency.
 - For "Suffix": Enter .jpg (or .png, .jpeg, etc., separated by | if you want multiple). This ensures your Lambda only triggers for actual image files, saving unnecessary invocations.
 - "Enable trigger": Make sure this is checked.
4. Finalize Trigger Addition: Click "Add".

Upload Image Processing Code (with Pillow) via ZIP

Lambda environments are minimal by default. For image manipulation, we need libraries like Pillow, which aren't built-in. This means we'll create a "deployment package" – a ZIP file containing our Lambda code and all its dependencies (Pillow in this case).

What our Lambda will do:

- Receive a trigger event from the S3 input bucket.
- Download the newly uploaded image.
- Resize it to 128x128 pixels.

- Upload the resized image to the S3 output bucket.

Required Files:

1. `lambda_function.py`: This is your core Lambda handler script.

```
from PIL import Image
import boto3
import os
import json
```

```
s3 = boto3.client('s3')
```

```
def lambda_handler(event, context):
    print("Lambda triggered with event:")
    print(json.dumps(event))
```

```
try:
```

```
    # Extract bucket and key from the S3 event
```

```
    input_bucket = event['Records'][0]['s3']['bucket']['name']
```

```
    input_key = event['Records'][0]['s3']['object']['key']
```

```
    # Get the output bucket name from environment variables
```

```
    # This is important for flexibility and security
```

```
    output_bucket = os.environ.get('OUTPUT_BUCKET')
```

```
    if not output_bucket:
```

```
        raise ValueError("OUTPUT_BUCKET environment variable is not set.")
```

```
    # Define the key for the processed image in the output bucket
```

```
    output_key = f"resized-{input_key}"
```

```
    print(f"Downloading from bucket: {input_bucket}, key: {input_key}")
```

```
    # Download the image to the Lambda's temporary storage
```

```
    s3.download_file(input_bucket, input_key, '/tmp/input.jpg')
```

```
    print("Opening and resizing image...")
```

```

# Open, resize, and save the image
img = Image.open('/tmp/input.jpg')
img = img.resize((128, 128)) # Resizing to 128x128
img.save('/tmp/output.jpg')

print(f"Uploading to bucket: {output_bucket}, key: {output_key}")
# Upload the processed image to the output bucket
s3.upload_file('/tmp/output.jpg', output_bucket, output_key)

print("Process completed successfully.")
return {"status": "Image resized and uploaded."}

except Exception as e:
    print(f"Error occurred: {str(e)}")
    # Re-raise the exception for CloudWatch logging
    raise e

```

Important Note: The code uses `os.environ['OUTPUT_BUCKET']`. This means we'll need to set an environment variable in your Lambda function configuration named `OUTPUT_BUCKET` with the value `image-processing-output`. We'll do this after uploading the code.

Step-by-Step to Create the Deployment Package using Docker:

Using Docker is the most reliable way to create a Lambda deployment package with specific dependencies like Pillow, ensuring compatibility with the Lambda execution environment.

1. **Create Project Directory:** Open your terminal or command prompt and run:


```
mkdir lambda-pillow
cd lambda-pillow
```
2. **Create Your `lambda_function.py`:** Inside the `lambda-pillow` directory, create a file named `lambda_function.py` and paste the Python code provided above into it.
3. **Create Dockerfile:** In the same `lambda-pillow` directory, create a file named `Dockerfile` (no extension) and paste the following content:

```
FROM python:3.12-slim
RUN apt-get update && apt-get install -y zip
WORKDIR /package
# Install Pillow here (NOT in /app)
RUN pip install Pillow -t .
# Now copy your lambda handler
COPY lambda_function.py .
CMD ["bash"]
```

4. Build the Package in Docker: Make sure you have Docker Desktop running. In your terminal, from within the lambda-pillow directory, run:

docker build -t lambda-packager .

This command builds a Docker image that contains Python 3.12, Pillow, and your Lambda code.

5. Zip Your Package: Now, we'll use the Docker image to create the ZIP file.
 - For Windows (using Command Prompt or PowerShell):
docker run --rm -v "%cd%:/output" lambda-packager bash -c "cd /package && zip -r9 /output/lambda-pillow.zip ."
 - For macOS/Linux (using Bash):
docker run --rm -v "\$(pwd):/output" lambda-packager bash -c "cd /package && zip -r9 /output/lambda-pillow.zip ."
6. This command runs the Docker container, zips up the Pillow library and your lambda_function.py into lambda-pillow.zip, and places it in your current lambda-pillow directory.
7. Check Contents (Optional but Recommended): You can inspect the contents of the ZIP file to ensure everything is there:
 - For macOS/Linux:
tar -tvf lambda-pillow.zip
 - For Windows (you might need a tool like 7-Zip or just open the .zip file directly): You can simply navigate to your lambda-pillow directory in File Explorer and open lambda-pillow.zip to confirm lambda_function.py and the PIL directory (from Pillow) are present.